



**High-Performance Data-Intensive
Computing Systems Laboratory**

National Research Platform Nautilus Research Cluster

MORENet Technical Summit

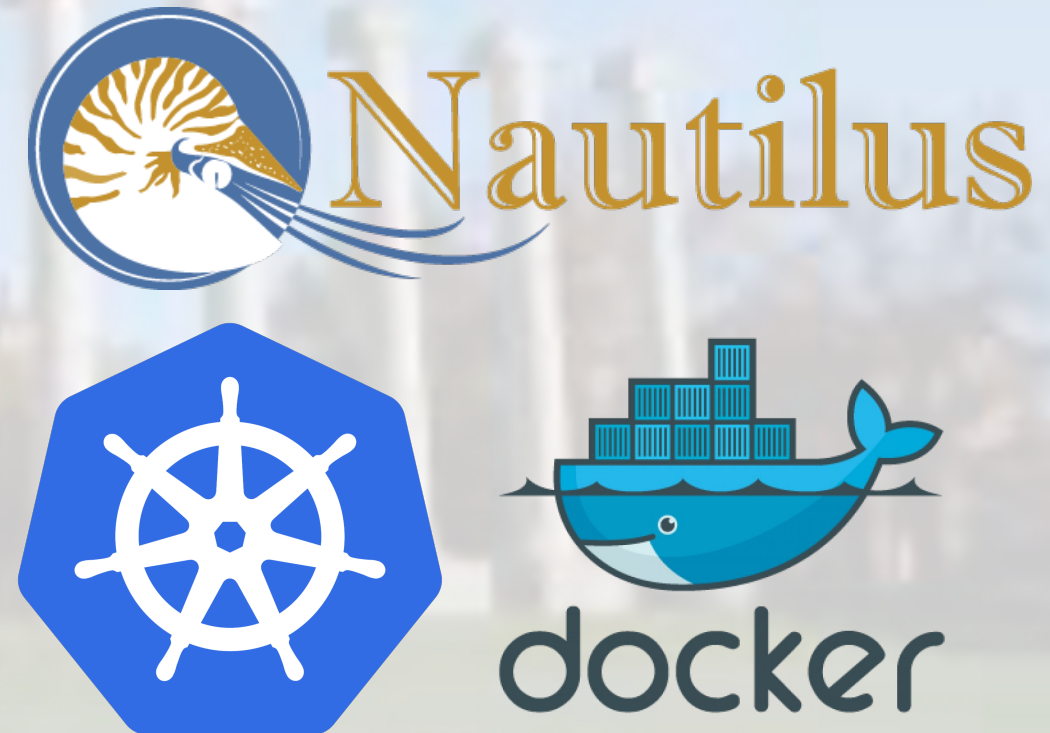
20 Feb 2023



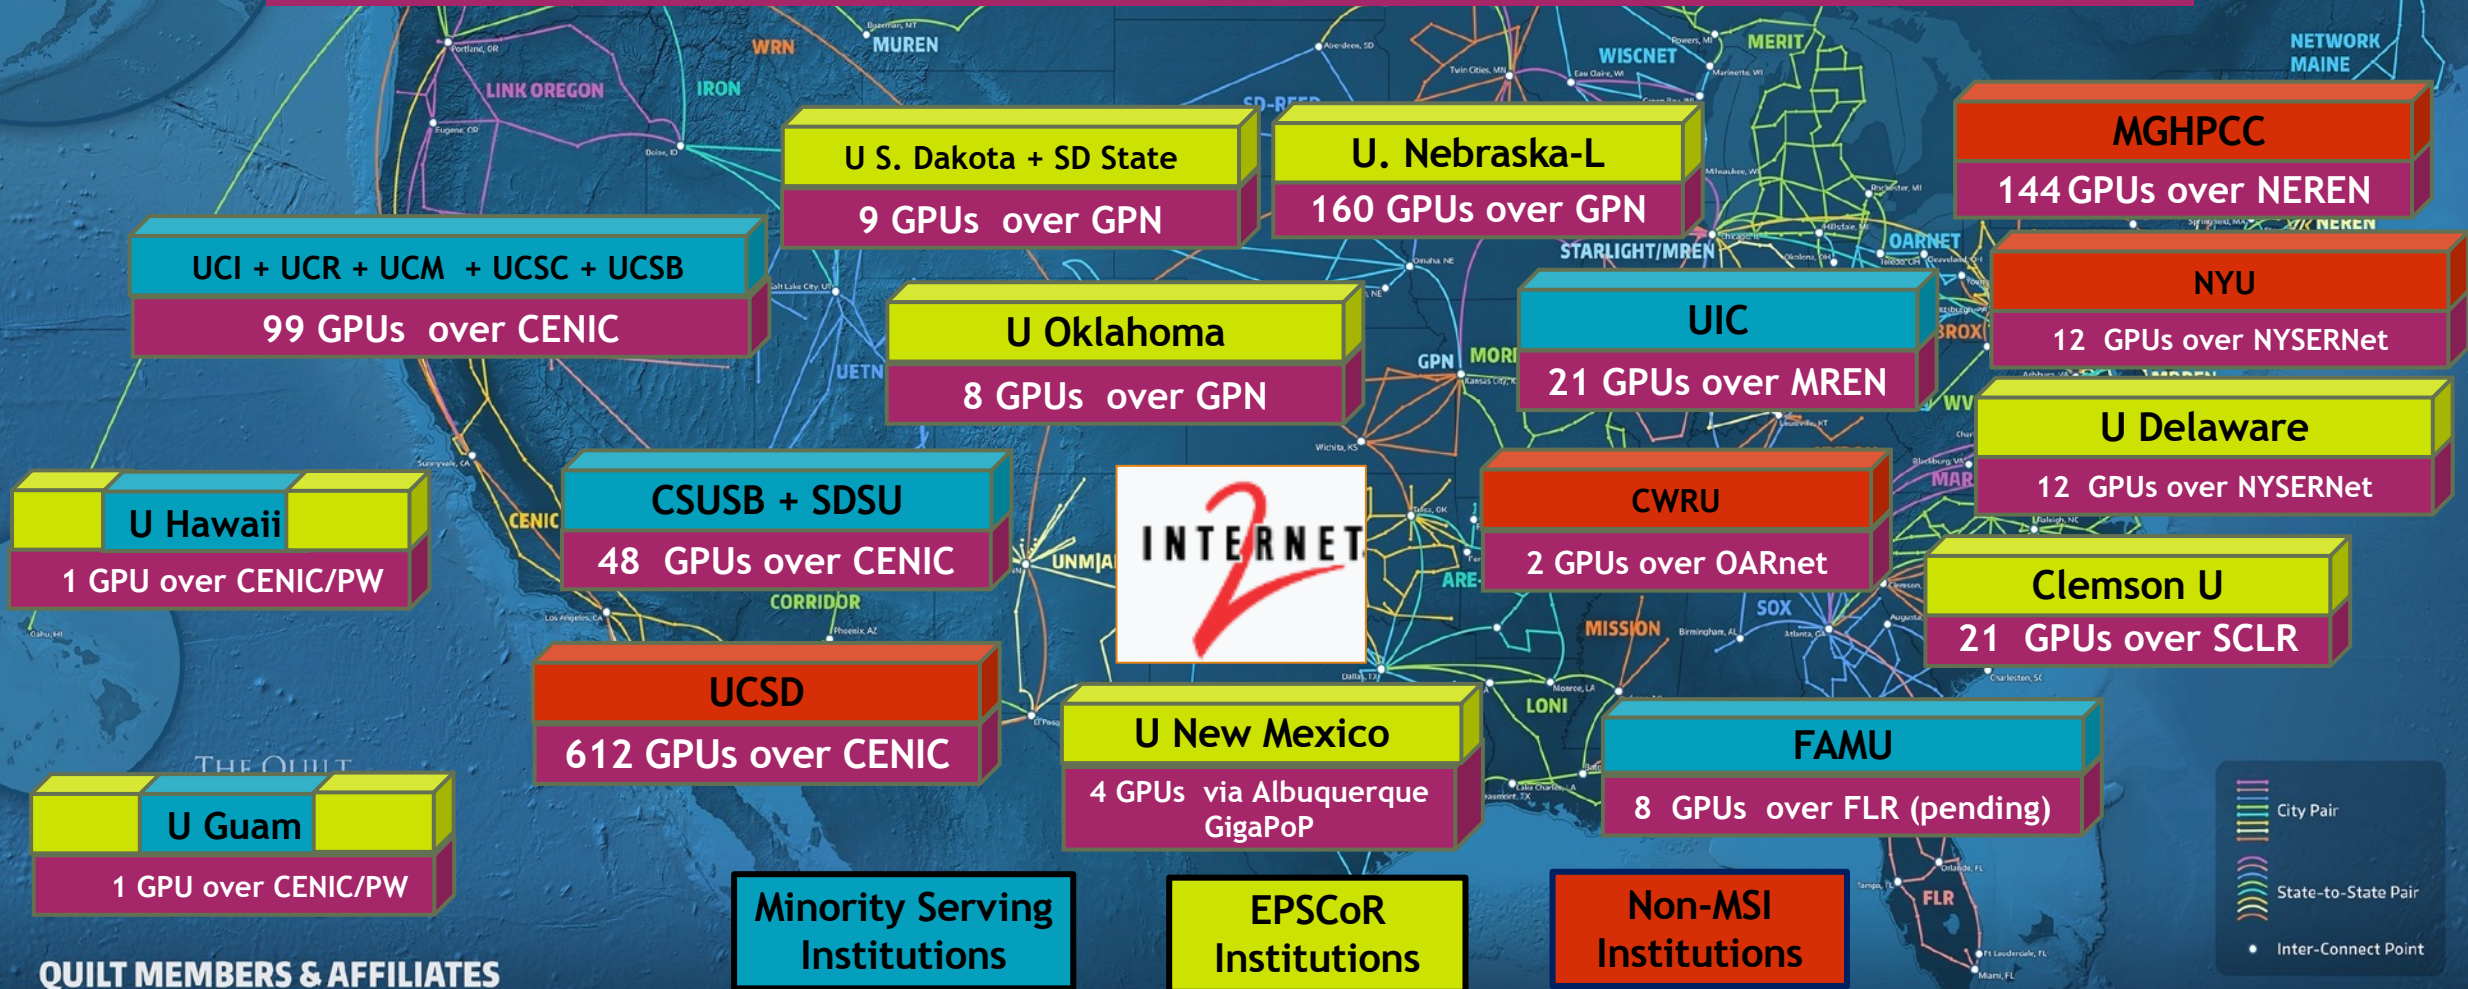
University of Missouri

NSF NRP Nautilus HyperCluster

- ▶ The NSF Nautilus HyperCluster is a Kubernetes cluster with vast resources that can be utilized for various research purposes:
 - ▶ Prototyping research code
 - ▶ S3 cloud storage for data and models
 - ▶ Accelerated small-scale research compute
 - ▶ Scaling research compute for large scale experimentation
- ▶ Resources Available:
 - ▶ CPU Cores: 14,462
 - ▶ RAM: 69 TB
 - ▶ NVIDIA GPUs: 1150



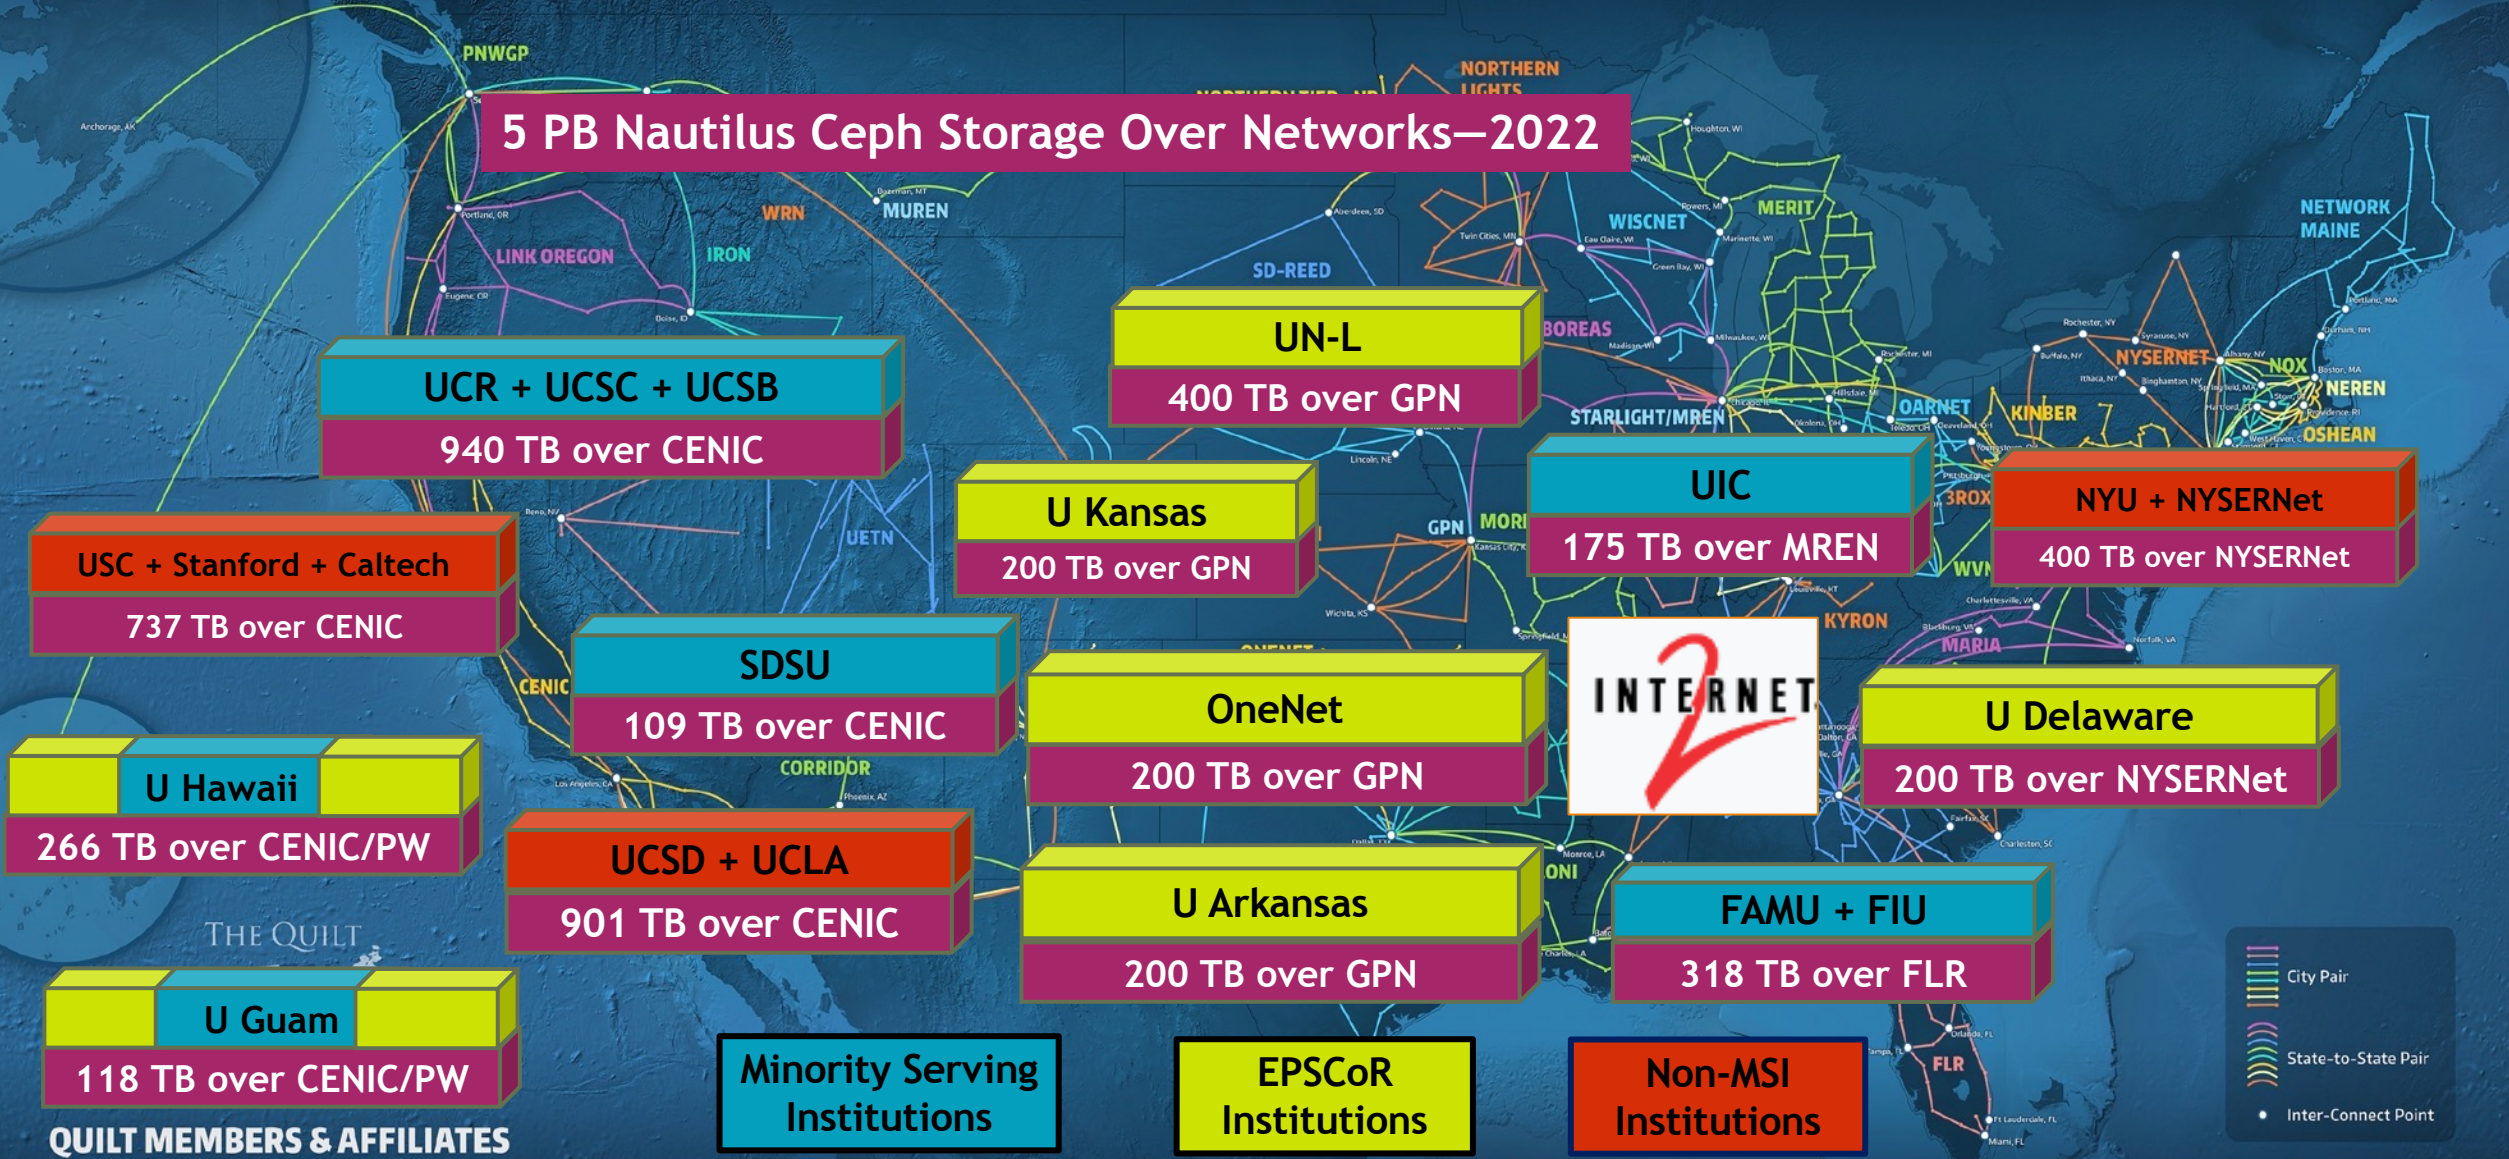
Nautilus ~1,100 GPUs Distributed over US Networks—Fall 2022



QUILT MEMBERS & AFFILIATES



5 PB Nautilus Ceph Storage Over Networks—2022



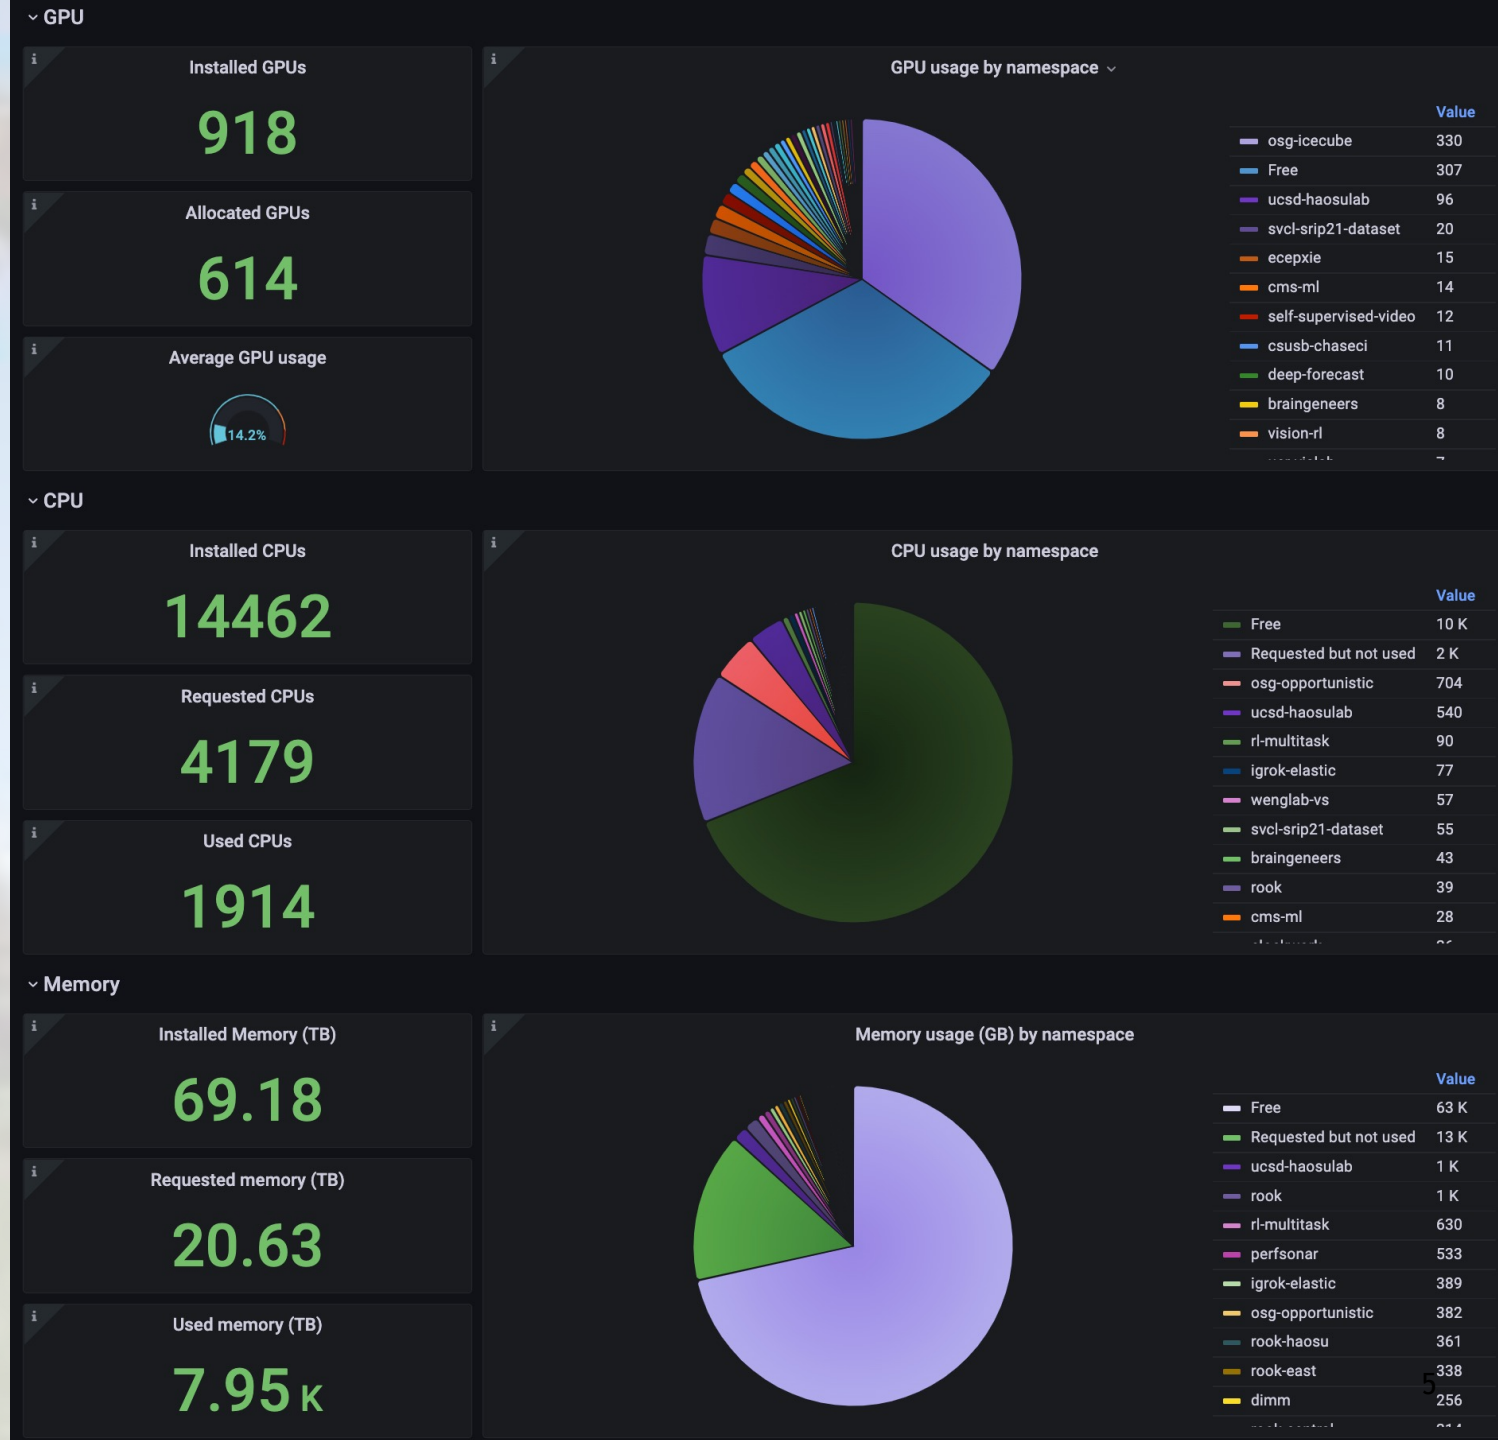
QUILT MEMBERS & AFFILIATES



Resource Dashboard

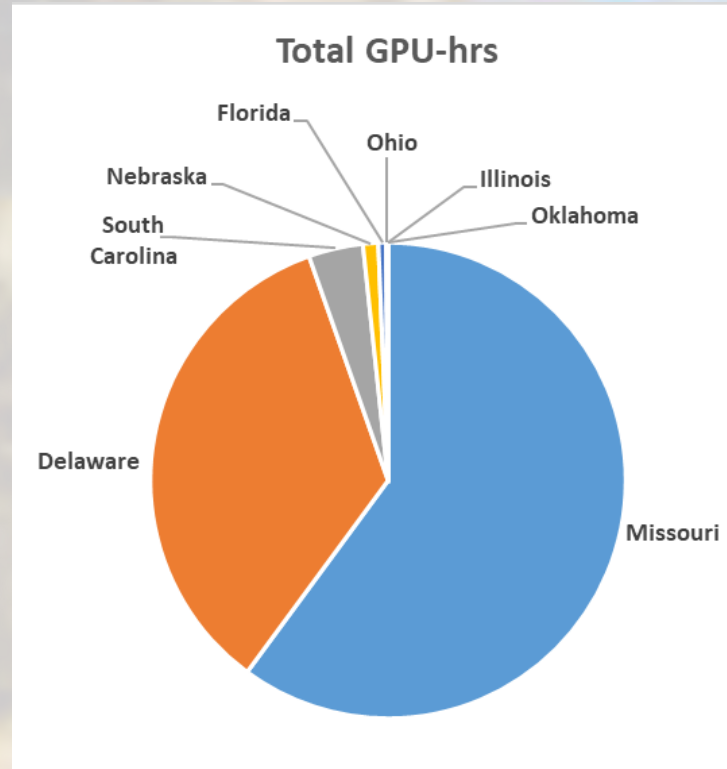


University of Missouri

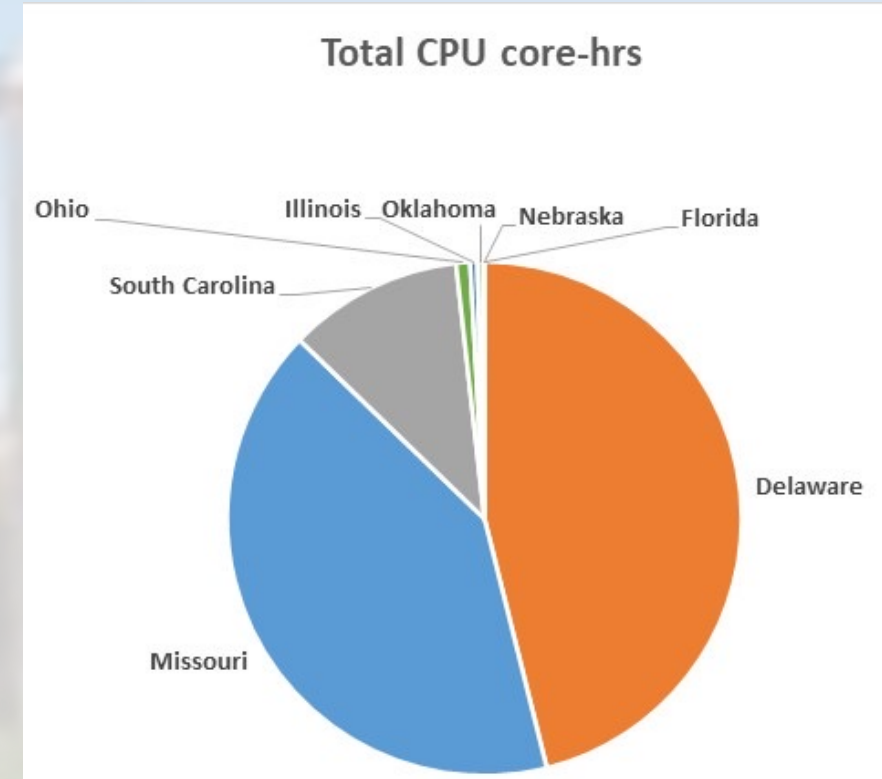


Non-California Nautilus PI Namespace 2021 Usage by State: “Big MO!”

Data/Plots provided by Larry Smarr (PI, National Research Platform & father of US Super Computing Centers)



17,217 GPU-hrs



28,088 CPU core-hrs

University of Missouri - Columbia:
42,000 GPU-hrs in 2022!



University of Missouri

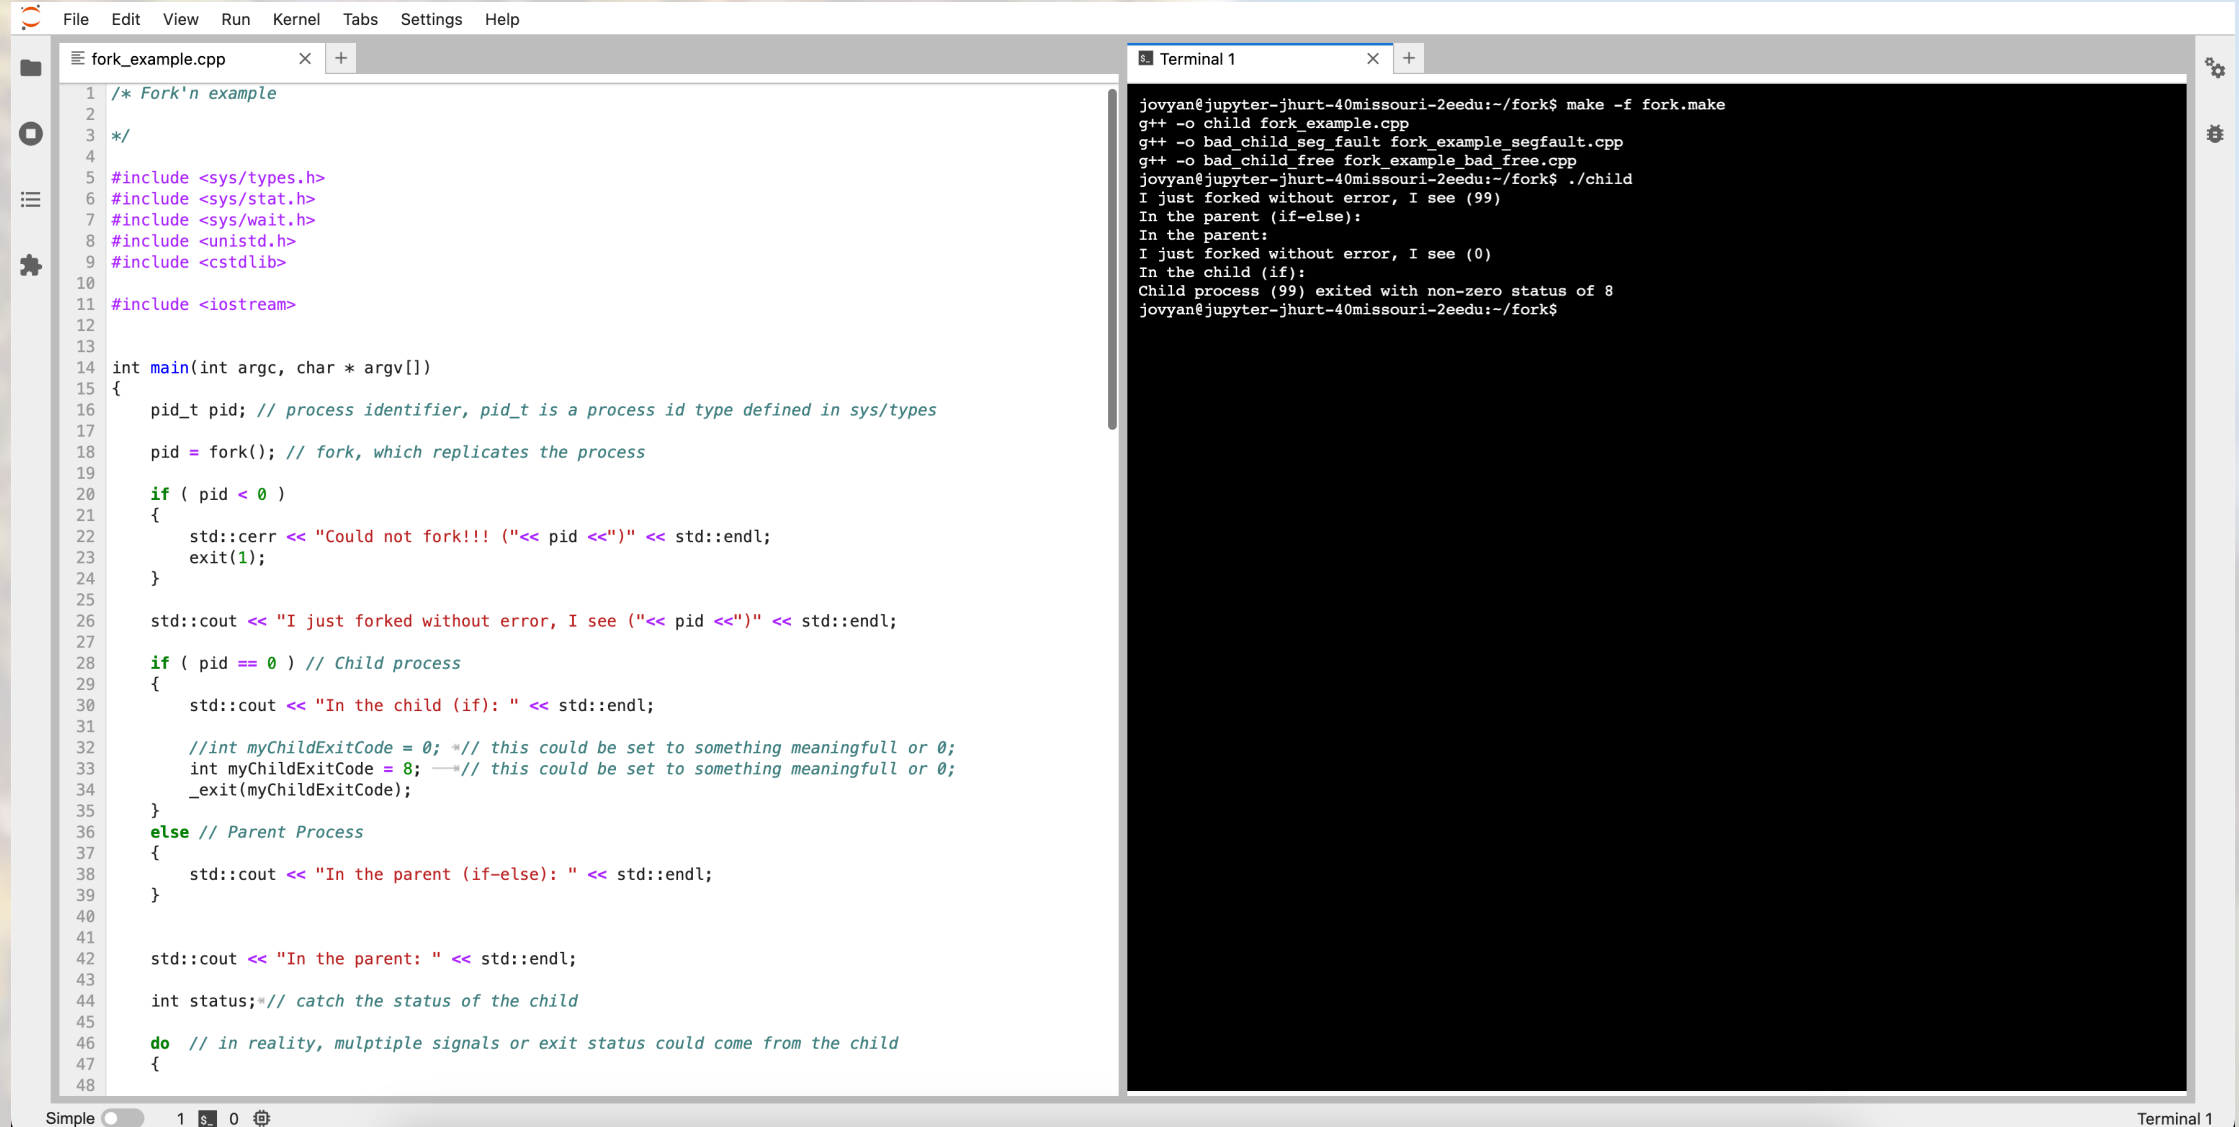
Grant Scott, UMC
Helped Organize the UMC PRP Usage

How MU is using Nautilus: Teaching & Research

- ▶ JupyterHub integration enables creation of an interactive learning environment for STEM Education:
 - ▶ Data Science
 - ▶ Mathematics
 - ▶ High Performance Computing
- ▶ Access to vast amount of RAM, CPU Cores, and NVIDIA GPUs can accelerate research in various fields:
 - ▶ Bioinformatics
 - ▶ Remote Sensing
 - ▶ Materials Science
 - ▶ Computer Vision
 - ▶ Machine Learning



HPC Interactive Learning Environment



The image shows a C++ IDE window titled 'fork_example.cpp' with the following code:

```
1  /* Fork'n example
2
3  */
4
5  #include <sys/types.h>
6  #include <sys/stat.h>
7  #include <sys/wait.h>
8  #include <unistd.h>
9  #include <cstdlib>
10
11 #include <iostream>
12
13
14 int main(int argc, char * argv[])
15 {
16     pid_t pid; // process identifier, pid_t is a process id type defined in sys/types
17
18     pid = fork(); // fork, which replicates the process
19
20     if ( pid < 0 )
21     {
22         std::cerr << "Could not fork!!! (" << pid << ")" << std::endl;
23         exit(1);
24     }
25
26     std::cout << "I just forked without error, I see (" << pid << ")" << std::endl;
27
28     if ( pid == 0 ) // Child process
29     {
30         std::cout << "In the child (if): " << std::endl;
31
32         //int myChildExitCode = 0; // this could be set to something meaningful or 0;
33         int myChildExitCode = 8; // this could be set to something meaningful or 0;
34         _exit(myChildExitCode);
35     }
36     else // Parent Process
37     {
38         std::cout << "In the parent (if-else): " << std::endl;
39     }
40
41     std::cout << "In the parent: " << std::endl;
42
43     int status; // catch the status of the child
44
45     do // in reality, multiple signals or exit status could come from the child
46     {
47
48
```

The terminal window titled 'Terminal 1' shows the following output:

```
jovyan@jupyter-jhurt-40missouri-2eetu:~/fork$ make -f fork.make
g++ -o child fork_example.cpp
g++ -o bad_child_seg_fault fork_example_segfault.cpp
g++ -o bad_child_free fork_example_bad_free.cpp
jovyan@jupyter-jhurt-40missouri-2eetu:~/fork$ ./child
I just forked without error, I see (99)
In the parent (if-else):
In the parent:
I just forked without error, I see (0)
In the child (if):
Child process (99) exited with non-zero status of 8
jovyan@jupyter-jhurt-40missouri-2eetu:~/fork$
```