Running JupyterHub in Kubernetes

MORENet Technical Summit

20 Feb 2023

University of Missouri

The Problem: Scalability & Reproducibility

- How do we ensure reliable portability of software developed on local development machines to other computational environments?
- How do we move code from local development on one machine to hundreds/thousands of machines?

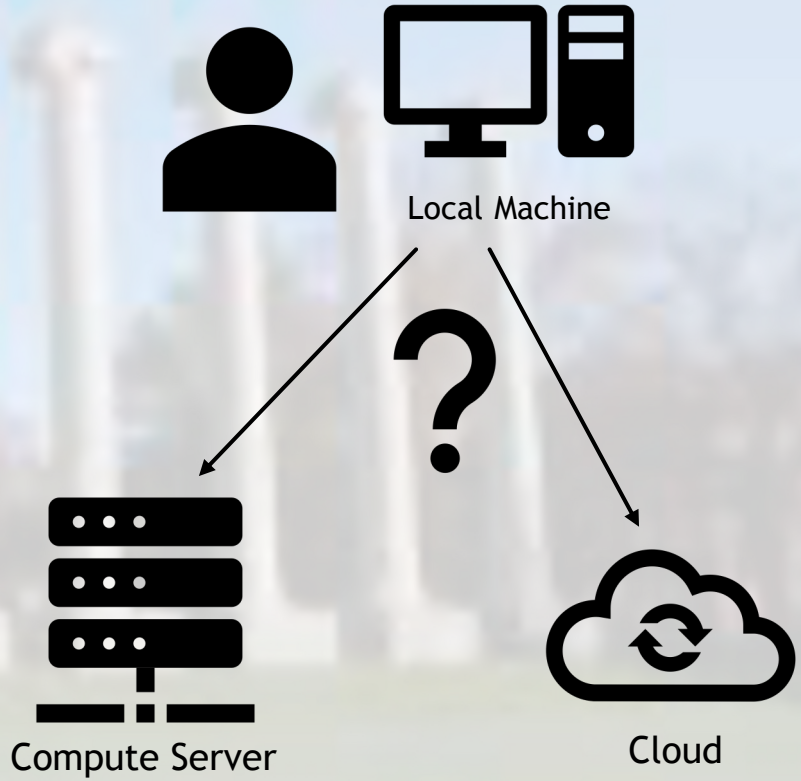

JupyterHub

- JupyterHub is a browser-based interactive environment that can be utilized for many different applications
  - Teaching
  - Research
- Using JupyterHub for STEM instruction provides an opportunity for hands-on, interactive education
- JupyterHub can authenticate with CILogin, meaning no additional authentication management is necessary

Running JupyterHub on Nautilus

1. Register CiLogon Application
2. Create Kubernetes Namespace
3. Request removal of Pod Limitations
4. Install Helm
5. Build Configuration YAML
6. Deploy JupyterHub

Step 1: Register CiLogon Application

- https://cilogon.org/oauth2/register
  - Callback URL: https://YOURNAME.nrp-nautilus.io/hub/oauth_callback
  - Client Type: Confidential
  - Scopes:
    - org.cilogon.userinfo
    - openid
    - Profile
    - Email
  - Refresh Tokens: No
- **Important: Keep the Client ID and Client Secret!**

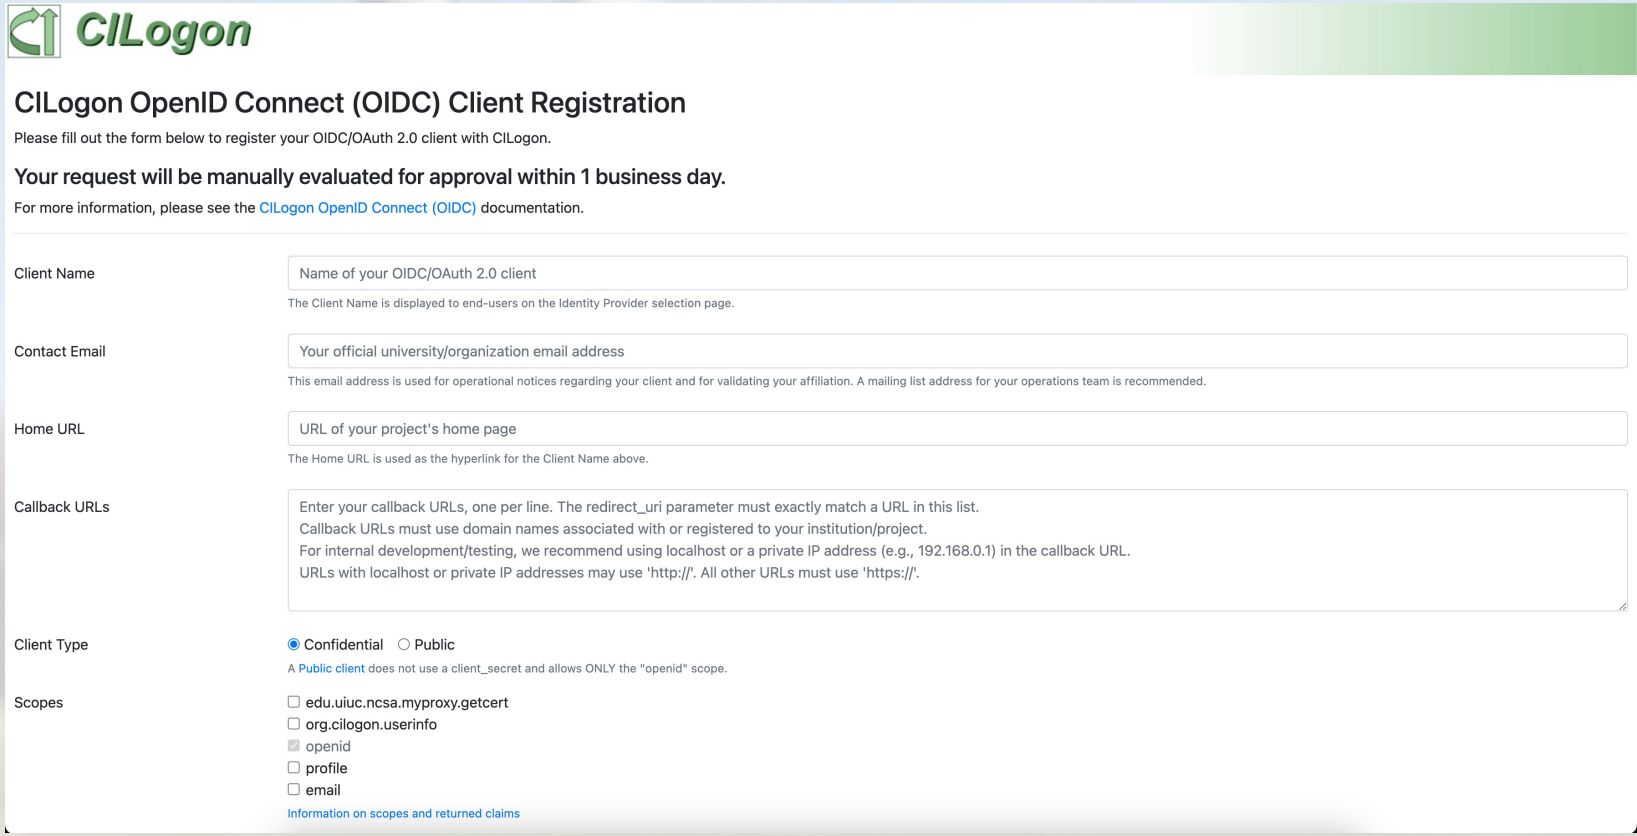

Step 2: Create a Kubernetes Namespace

- Recall: **Namespaces** provide a way for K8s to partition cluster resources across multiple or many users in an exclusive way.
- All of the instances of Jupyter Hub running will be in a separate pod in a namespace
- To create the namespace on the Nautilus cluster, a cluster administrator will need to go to the following URL: https://portal.nrp-nautilus.io/profileN

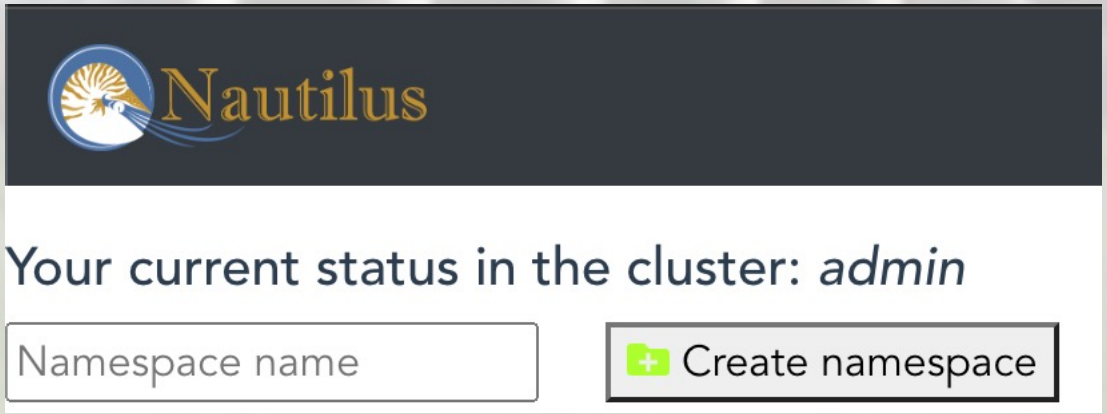

Step 3: Request Removal of Pod Restrictions

- By default, pods on the Nautilus cluster can only run for 6 hours
- If you have users that may use their Jupyter Hub environment for more than 6 hours, their resources will be released at the 6 hour mark
- You can request that pod restrictions be removed on your namespace by going to the [Matrix] chat for PRP Nautilus here: https://element.nrp-nautilus.io/#/welcome
- Once you have access to the server, send a message in **Nautilus Support** channel containing the name of the namespace you are using for Jupyter Hub

Step 4: Install Helm

- Helm is a package manager for Kubernetes, and enables software deployment on K8s Clusters
- We will use helm to manage our JupyterHub installation
- To download and install: https://helm.sh/docs/intro/install

Step 5: Build Configuration YAML File

- Helm is configured, similarly to K8s, with YAML files
- We need to build a YAML config file to tell Helm how we want our JupyterHub application configured
- We can download the template for JupyterHub on Nautilus here:
  - https://ucsd-prp.gitlab.io/userdocs/jupyter/values.yaml
- Fields to Update in the Template:
  - `secret_token` should be updated to the output of this command: `openssl rand -hex 32`
  - `client_id` and `client_secret` should be updated to match what you received when you registered you CiLogon client
  - `admin_users` should be set to who you'd like to have administrative access inside Jupyter Hub
  - `oauth_callback_url` should updated to: https://YOURNAME.nrp-nautilus.io/hub/oauth_callback
  - `ingress.hosts` should be updated to: https://YOURNAME.nrp-nautilus.io/

Step 6: Deploy JupyterHub with Helm

- Add the Jupyter Hub repository:

Bash

```
$ helm repo add jupyterhub https://jupyterhub.github.io/helm-chart/ &&
  helm repo update
```

- Deploy the application:

Bash

```
$ helm upgrade --cleanup-on-fail --install jhub jupyterhub/jupyterhub
  --namespace <namespace> --version=1.2.0 --values config.yaml
```

  - where `config.yaml` is the configuration file you setup in Step 5 and `<namespace>` is the namespace you created in Step 2.